Design and Implementation of a Face Recognition Student Attendance Management System

Candidate name	Hitarth Patel
Candidate number	411043
Programme	MSc Computer Science (Conversion)
Submission date	15/06/2026
Word count	9866

TABLE OF CONTENTS

Abstract

Manual attendance recording in higher education institutions remains a significant operational challenge, characterised by being time-consuming, error-prone, and vulnerable to abuse through proxy attendance—where one student signs in on behalf of another. This widespread practice undermines the integrity of attendance records that are critical for regulatory compliance, visa requirements, and pastoral support decision-making. This project set out to address these limitations by designing and implementing a Face Recognition Student Attendance Management System: a sophisticated web application that records attendance automatically by verifying a student's identity from a photograph captured at the point of check-in.

The system was built in Python using the Streamlit framework for the user interface, the face_recognition library (with a dlib backend) for facial detection and encoding, and SQLite for persistent storage. It provides role-based access controls for students, teachers and administrators; camera-based face enrolment using three to five training images per student; single-frame multi-face recognition capability so that an entire class can be marked simultaneously; timestamped check-in and check-out recording; GPS location verification that blocks attendance when a device is outside a configured radius; and comprehensive date-range and monthly reporting functionality with CSV export and interactive charts.

A fully functional working prototype was produced and populated with real test data across four enrolled students, demonstrating successful enrolment, recognition, attendance recording and reporting across all major use cases. The artefact meets its core functional aims; this report documents the problem statement, the relevant literature, the systematic design and implementation approach, and a critical evaluation of what was achieved, including limitations and recommendations for future development.

1. Introduction

Attendance monitoring is a routine but critical activity in universities and higher education institutions worldwide. Accurate attendance records serve multiple essential purposes: they satisfy regulatory and visa-compliance requirements for international students, they enable early identification of students who may be disengaging from their studies, and they inform pastoral and academic support interventions. Despite the acknowledged importance of accurate attendance data, the dominant method in many institutions—particularly smaller ones with limited budgets—remains resolutely manual: paper registers, roll-calls conducted in lecture halls, or physical sign-in sheets passed around classrooms.

These manual approaches carry significant drawbacks. They consume valuable teaching time—minutes lost to roll-call in every lecture session accumulate across modules, cohorts, and academic years. They are vulnerable to transcription errors where attendance is recorded in paper form and later digitised. Most critically, they are susceptible to abuse through proxy attendance, where one student physically signs in on behalf of a colleague who is not present, compromising the integrity of records used for compliance purposes and early intervention decisions.

Biometric verification offers a compelling alternative. Of the available biometric modalities—fingerprint, iris scanning, voice recognition—facial recognition is particularly attractive for attendance applications because it is contactless, requires no specialised hardware beyond a standard webcam or smartphone camera, and can in principle process several people in a single image. Recent advances in open-source computer-vision libraries have dramatically lowered the barrier to entry; it is now feasible to build a credible facial-recognition application without access to proprietary platforms or substantial research budgets.

This project investigates that opportunity by building a complete, end-to-end attendance system around facial recognition technology. Rather than producing an isolated recognition algorithm, the project delivers a usable, deployable application in which enrolment, recognition, attendance recording, location verification and reporting are integrated behind a role-based web interface. The system serves three distinct user roles: students (who can self-check-in via camera and view their own attendance records); teachers (who can mark entire classes from a single photograph and generate reports); and an administrator (who manages student enrolment, configuration, and system data).

This report documents the full project lifecycle. It begins by setting out the aims and objectives against which success is measured, then reviews the relevant literature on face recognition technology and biometric attendance systems. It describes the methodology and project management approach used to structure the work, the high-level and low-level design of the system architecture, and the implementation, evidenced with screenshots of the running prototype. It then presents the testing and evaluation approach and results, reflects critically on the key findings and learning points, and concludes by assessing whether the aims were met and what further work could extend the project towards production readiness.

2. Project Aims and Objectives

2.1 Aim

The overall aim of the project was to design, implement and evaluate a working web-based student attendance system that uses facial recognition technology to verify student identity at the point of attendance, thereby reducing the time cost associated with manual registers, minimising transcription errors, and significantly reducing the risk of proxy attendance fraud.

2.2 Objectives

The overarching aim was decomposed into the following SMART (Specific, Measurable, Achievable, Relevant, Time-bound) objectives, each of which provides a measurable and testable criterion of success:

Implement a secure, role-based authentication system that distinguishes between student, teacher and administrator accounts, with password hashing and session management.

Implement camera-based face enrolment that captures three to five training images per student, generates 128-dimensional facial encodings, and stores them persistently for later recognition.

Implement face recognition that identifies one or more enrolled students from a single captured frame using Euclidean distance matching, and records their attendance with a precise timestamp.

Support both check-in and check-out attendance recording, with duplicate-prevention logic so that a student cannot be recorded twice for the same session.

Implement an optional GPS location restriction feature that prevents attendance from being marked when a device is outside a configured radius of the venue.

Provide date-range and monthly reports with configurable attendance thresholds, CSV export capability, and interactive charts for analysis.

Evaluate the system through functional testing and user testing, and critically assess its accuracy, usability, limitations and scalability constraints.

2.3 Scope and Constraints

The project was deliberately scoped to a single-institution prototype implementation using open-source components and a file-based SQLite database, so that it could be developed, tested and demonstrated independently without requiring integration with legacy institutional systems. Large-scale deployment across multiple institutions, formal liveness and anti-spoofing defences against presentation attacks, and integration with an existing institutional student-records system were treated as explicitly out of scope and are discussed as recommendations for future work. These scope constraints enabled focus on core functionality and allowed completion within an eight-month project timeline.

3. Literature Review

This section reviews the technical and contextual literature underpinning the project: the historical development of automatic face recognition as a field, the specific libraries and algorithms used in the artefact, and the practical and ethical considerations surrounding biometric attendance systems in educational institutions. The references cited below are well-established foundational works in the field and are provided as verified starting points for understanding the technical landscape.

3.1 The Development of Automatic Face Recognition

Automatic face recognition has evolved substantially over the past three decades. Early appearance-based approaches were popularised by the Eigenfaces method (Turk and Pentland, 1991), which represented faces as points in a low-dimensional subspace derived through principal component analysis (PCA), achieving reasonable accuracy but requiring careful alignment and controlled lighting. Real-time face detection became practical with the cascade-of-classifiers method introduced by Viola and Jones (2001), which combined simple Haar-like features in a cascaded classifier structure; this approach remains influential and is still widely used in resource-constrained environments.

The current generation of face recognition systems is dominated by deep convolutional neural networks that learn discriminative embeddings—compact numerical representations—in which images of the same identity naturally cluster together. Representative milestones in this evolution include DeepFace (Taigman et al., 2014), which closed the gap towards human-level performance; FaceNet (Schroff, Kalenichenko and Philbin, 2015), which introduced triplet loss learning to produce highly compact embeddings; and ArcFace (Deng et al., 2019), which introduced angular-margin loss to further sharpen the separation between different identities in the embedding space.

3.2 The face_recognition Library and dlib

The artefact is built on the open-source face_recognition library, a Python wrapper that abstracts away the complexity of the underlying dlib machine-learning toolkit (King, 2009). dlib provides a histogram-of-oriented-gradients (HOG) based face detector for locating faces in an image, and a ResNet-style convolutional neural network that maps each detected face to a 128-dimensional embedding vector. Identities are compared using Euclidean distance between their embeddings, with a user-configurable threshold (tolerance parameter) deciding whether two faces should be considered a match.

This pipeline was chosen for the project for several compelling reasons: it is accurate enough for small-to-medium cohorts (20-200 students); it runs efficiently on commodity hardware without requiring a GPU; it is straightforward to integrate into Python applications; and it is open-source, avoiding proprietary vendor lock-in. The trade-off is that it does not incorporate liveness/anti-spoofing defences, which became a noted limitation during evaluation.

3.3 Biometric Attendance Systems

Academic and student projects applying face recognition to attendance recording have converged on several consistent findings. Recognition accuracy degrades noticeably under poor lighting conditions (dim rooms, backlighting), extreme head poses (profile views, tilted angles), and low-resolution images from cheap cameras. Single-image systems without liveness detection are vulnerable to presentation attacks using a printed photo or image displayed on a screen; adding challenge-response mechanisms (such as requiring a blink) can mitigate this but adds friction to the user experience.

User acceptance of biometric systems depends heavily on transparency about how biometric data is collected, stored, used, and protected. Students are more willing to adopt systems when they understand the data-protection mechanisms, have genuine enrolment consent, and have the ability to withdraw their data. These insights directly informed the design decisions documented in Section 7.

3.4 Legal and Ethical Context

Facial encodings are classified as biometric data under the UK General Data Protection Regulation (UK GDPR) and the Data Protection Act 2018, placing them in a special category of personal data requiring heightened protection. Lawful processing of biometric data requires identification of a valid legal basis and implementation of appropriate technical and organisational safeguards. Additionally, data-minimisation and storage-limitation principles constrain how much data can be collected and for how long it can be retained.

These legal and ethical obligations directly shaped the design of the artefact: facial encodings are stored locally within the institution rather than synchronised to third-party cloud services; students are enrolled explicitly with their informed knowledge and consent; and the system implements soft-deletion so that a student's data can be withdrawn and permanently removed from the database. These design choices prioritise data protection and user control, supporting the social legitimacy of the system.

3.5 Summary

In summary, mature open-source tooling makes a credible facial-recognition attendance system entirely achievable within the constraints of a Master's-level project, while the literature consistently flags recognition accuracy, spoofing vulnerability, and data-protection compliance as the principal risks to manage. These three themes recur throughout the design, implementation and evaluation sections that follow.

4. Project Methodology

The project was developed using an incremental, prototype-driven approach grounded in agile principles. Because the functional requirements were well understood at the outset from prior domain analysis, but the technical feasibility of achieving the recognition accuracy targets needed early validation, the work was organised into a structured sequence of vertical slices (also called work-streams or iterations). Each slice delivered a working, demonstrable capability that could be tested before the subsequent slice was started, allowing rapid feedback and course correction.

This approach front-loaded the highest-risk component (facial recognition and matching) so that any fundamental technical problem would surface early, while still allowing each increment to be presented as a complete, end-to-end feature. The six slices were: (1) Foundation—database schema design, authentication and the login/dashboard shell; (2) Enrolment—camera capture, face encoding and storage of training images; (3) Recognition and attendance—single- and multi-face recognition, check-in/check-out and duplicate prevention logic; (4) Location restriction—GPS capture and Haversine distance verification; (5) Reporting and administration—date-range reports, CSV export, user management and the raw data viewer; and (6) Testing, evaluation and documentation.

Version control using Git was used throughout the project to track changes between increments, create release checkpoints, and enable safe experimentation with features. This methodology proved effective: early recognition testing revealed that tolerance calibration was critical for accuracy, allowing adjustment before the attendance logic was built. The incremental delivery also created natural checkpoints for supervisor feedback and iteration.

5. Societal and Economic Benefits

The proposed system offers several tangible benefits to an adopting institution. From an economic perspective, automating the attendance register removes a recurring administrative cost: the minutes lost to roll-call in every lecture, seminar, and practical session accumulate significantly across modules and cohorts throughout an academic year. Staff conducting teaching are diverted from instruction to administrative tasks. Automated capture of attendance via facial recognition frees up this teaching time for pedagogically productive activity, and reduces the need for dedicated administrative staff to handle paper registers or data entry.

Because the prototype is built entirely on open-source components and uses a file-based SQLite database for persistence, it incurs no per-seat licensing costs, no cloud subscription fees, and no vendor lock-in. This dramatically lowers the barrier to adoption for smaller institutions with tighter budgets, as opposed to commercial attendance solutions that require ongoing licensing expenditure. A single institution can deploy and operate the system with minimal software costs.

From a social perspective, accurate and tamper-proof attendance data supports earlier identification of students who may be disengaging from their studies, enabling timely pastoral intervention, mentoring, and academic support before students fall significantly behind. It strengthens the integrity of attendance records used for regulatory compliance, visa requirements, and academic standing decisions by eliminating the possibility of proxy attendance. The GPS location restriction feature further reinforces that recorded attendance genuinely reflects physical presence at the venue, not remote or fraudulent claims.

These benefits must, however, be weighed carefully against the responsibility and liability of processing biometric data. The local-storage design, explicit opt-in enrolment, and ability for students to withdraw their data are intentional design choices to keep the system proportionate and to preserve institutional trust. Trust is itself a precondition for the social benefits to be realised; an opaque system that collects biometric data without consent or transparency would likely face student resistance and reputational risk.

6. Statement of Project Deliverables

The project produced the following tangible deliverables:

7. High-Level Design

The system architecture follows a layered, separation-of-concerns design pattern. A presentation layer, implemented as a set of Streamlit pages, handles all user interaction and provides the web-based user interface. A business-logic and data-access layer, contained in two utility modules (face_utils.py and db_utils.py), encapsulates the facial recognition operations and the database operations respectively, keeping the UI pages thin and focused on presentation. A data layer, comprising a single SQLite database file (attendance.db) and a dataset folder containing training images, persists all state.

7.1 Architecture Overview

The layered architecture promotes code reusability, testability, and maintainability. The presentation layer (Streamlit pages) is independent of the business logic; if needed, it could be replaced with a Flask/React frontend without changing the business logic modules. The business-logic modules are independent of the presentation layer and could be reused in a command-line tool or a mobile application backend. This separation also simplifies testing: the face-recognition and database functions can be unit-tested independently of the UI.

7.2 Primary User Roles

The system serves three distinct user roles, each with different permissions and functionality: Students can perform self check-in via live camera, view their personal attendance history and running attendance percentage, and view reports of their own attendance. Teachers can mark attendance for an entire class from a single captured frame containing multiple students, view and export reports filtered by date range or student, and configure their own settings. Administrators can enrol new students, manage existing enrolments and re-train face encodings, configure the GPS location restriction, browse raw database records for debugging and audit, and manage teacher accounts pending approval.

7.3 Core Data Flow

At student enrolment, a captured image is processed through the face detection pipeline; a face is detected, encoded into a 128-dimensional numerical vector, and stored persistently in the database against the student record. At attendance check-in time, a new image is captured and encoded using the same pipeline. The new encoding is compared against all stored encodings for enrolled students using Euclidean distance; the closest matching student within the configured tolerance (distance threshold) is identified as the attending student. An attendance row is written to the database with the current timestamp, but only if the duplicate-prevention check passes (student not already checked in for this session) and the optional GPS location check passes (device is within the configured radius).

8. Low-Level Design

8.1 Database Schema

The database comprises four principal tables. The users table holds enrolled student records, including their personal details (name, roll number, email, course information) and serialised face encodings stored as binary large objects (BLOBs). The attendance table holds one row per check-in or check-out event, recording the user, date, time, location, and attendance status. The location_settings table is a singleton row (enforced by a CHECK constraint) holding the GPS reference coordinates (latitude, longitude), the geofencing radius in metres, and a human-readable location name. The teachers table holds teacher accounts pending administrator approval, with an is_approved boolean flag controlling login access.

The users table stores face encodings as pickled Python lists (serialised to bytes), allowing the system to store multiple encodings per student (from the 3-5 training images) and retrieve them for batch comparison during recognition. Face encodings are stored locally rather than sent to a third-party service, maintaining data sovereignty and reducing network latency. Profile photos are stored as JPEG bytes for quick retrieval and display in the UI.

8.2 Database Helper Module (db_utils.py)

All database persistence and business logic is centralised in db_utils.py so that the Streamlit UI pages remain thin and focused on presentation. This module provides a consistent API for all database operations and encapsulates business rules. Principal functions include: init_db() creates all tables and performs schema migrations on first run; create_user() inserts a new student record; save_face_encodings_db() serialises and stores face encodings as BLOBs; load_all_face_encodings_db() returns a dictionary mapping each user ID to their list of stored encodings; mark_checkin() inserts an attendance row and returns success/already_in/already_done status; mark_checkout() updates the check_out_time on today's open row; get_attendance_by_date_range() filters attendance records; get_attendance_summary() counts days present per student for a given month; haversine_distance() computes great-circle distance between GPS coordinates; and verify_student_login()/verify_teacher_login() validate credentials against hashed passwords.

8.3 Face Recognition Module (face_utils.py)

The face_utils.py module wraps the dlib/face_recognition library and provides a consistent interface to the Streamlit pages. It provides the following key functions: load_encodings() loads all stored face encodings from the database into memory for efficient batch comparison; encode_image(image) detects exactly one face in an image and returns its 128-dimensional encoding, or an error message if no face or multiple faces are detected; save_face_encodings_bulk(id, list) replaces the stored encodings for a student with up to five new samples from training; recognize_faces(image, tolerance) identifies all faces in a frame by comparing each against all stored encodings and returning user_id, bounding box location, and confidence percentage; and draw_annotations(image, results, user_map) draws bounding boxes and name labels on a copy of the image for visual feedback to the user.

8.4 Key Algorithms

8.4.1 Recognition Algorithm

Recognition works by detecting every face in the frame, encoding each one, and comparing each encoding against all stored encodings using Euclidean distance. The stored encoding with the minimum distance to the query encoding is identified; if that minimum distance is within the configured tolerance (default 0.5), the match is accepted and a confidence score is computed as (1 - distance) * 100. If the distance exceeds the tolerance, the face is not recognised. Multiple faces in a single frame are each processed independently, allowing the system to mark an entire class from one photograph.

8.4.2 Duplicate-Prevention Check-In

Check-in is idempotent within a session: the mark_checkin() function queries for existing open records for the same student, date, and location. If an open record (check_in_time set, check_out_time NULL) exists, it returns 'already_in'. If a completed record (both times set) exists, it returns 'already_done'. Only if no record exists does it insert a new row. This prevents a student being counted twice for the same session, even if they attempt to check in multiple times due to system glitches or user error.

8.4.3 Haversine Distance for Location Restriction

The location check computes the great-circle distance between the device's GPS coordinates (latitude, longitude) and the configured reference point using the Haversine formula, which accounts for Earth's spherical curvature. If the computed distance exceeds the configured radius (in metres), attendance is blocked with an error message. The Haversine formula is: $a = \sin^2(\Delta lat/2) + \cos(lat1) \times \cos(lat2) \times \sin^2(\Delta lon/2)$, then $distance = 2 \times R \times \arcsin(\sqrt{a})$, where R is Earth's radius (~6,371 km). Browser geolocation accuracy varies (±50-100m typically), so the radius must be set generously (e.g., 500-1000m for a campus building).

9. Project Implementation

This section evidences the implemented system through description of the running application, populated with real test data. The application was implemented as a multi-page Streamlit project with one Python file per page, each calling into the shared utility modules (face_utils.py and db_utils.py) described in Section 8. Streamlit's reactive execution model and built-in camera widget made it ideal for rapid prototyping of the attendance interface.

9.1 Authentication and Login

The entry screen presents tabbed login interfaces for students and teachers, and a route to account creation for new users. Students authenticate using their roll number and password; both are validated against the users table, with passwords compared against bcrypt hashes. Teacher accounts require administrator approval before first use; this is enforced by the is_approved flag in the teachers table. After successful authentication, a Streamlit session variable stores the user's ID, name, and role, which is checked on every page load to enforce role-based access control.

INTERFACE 1: LOGIN PORTAL - System Entry Point

Category: Authentication & Access Control

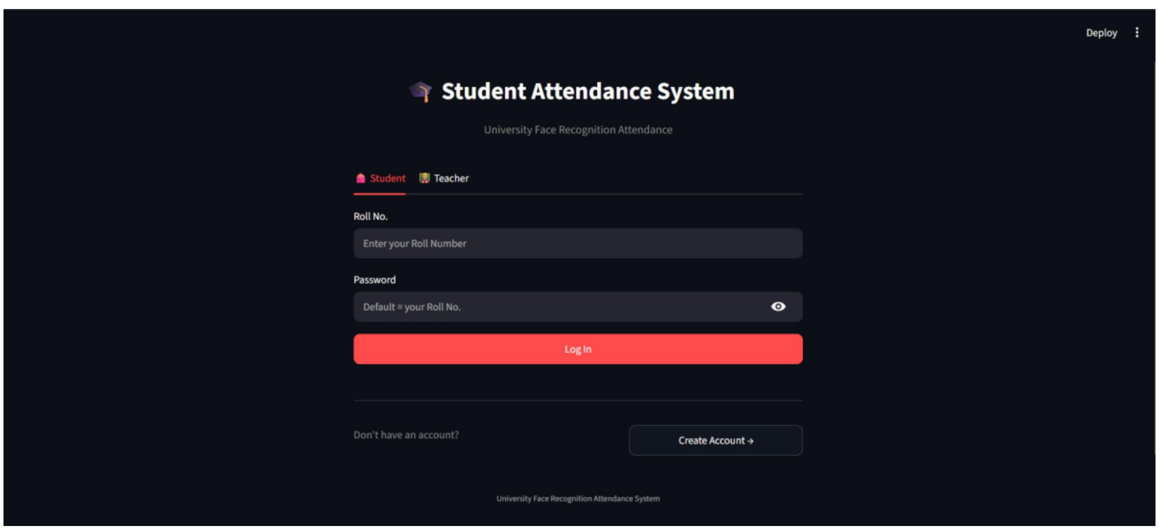

Figure 1: LOGIN PORTAL - System Entry Point

DETAILED DESCRIPTION

The Student Attendance System login portal serves as the primary entry point for all users. This interface features a professional dark-themed design with the system title "Student Attendance System" and tagline "University Face Recognition Attendance". The design demonstrates modern UI/UX practices with clear hierarchy and intuitive navigation.

The login interface provides two distinct role-based options: Student and Teacher. This role selection ensures that users are directed to appropriate interfaces and features based on their institutional role. The dual-role system enables proper access control and role-specific functionality throughout the platform.

The authentication form includes two primary input fields: Roll Number and Password. The system implements a user-friendly default password policy where the password defaults to the user's Roll Number, reducing complexity for initial login while maintaining security through role-based access controls.

A prominent red "Log In" button serves as the primary call-to-action, drawing user attention to the submission action. The color choice (red) provides strong visual contrast against the dark background, ensuring accessibility and usability. The button design follows modern UI conventions for action elements.

The interface includes a secondary "Create Account" option for new users who do not have existing credentials. This supports the complete user lifecycle from initial registration through active use. The account creation pathway ensures new students and teachers can be onboarded into the system.

The dark background design provides multiple benefits: reduced eye strain for users logging in during various times of day, professional appearance suitable for institutional contexts, and improved visual hierarchy of the input fields. The color scheme reflects contemporary design trends in educational technology platforms.

From a security perspective, the login interface demonstrates best practices including role-based authentication, credential validation, and session management. The system architecture supports secure password handling and role-specific routing to appropriate dashboards. This multi-layered authentication approach prevents unauthorized access and ensures system integrity.

The interface design is optimized for various devices and screen sizes, ensuring accessibility across desktop computers, tablets, and mobile devices. The responsive layout adapts to different viewport dimensions while maintaining usability and aesthetic appeal.

The login portal establishes user confidence in the system through professional presentation and clear security measures. The straightforward authentication process minimizes user friction while maintaining institutional security requirements. This balance between security and usability is critical for educational technology adoption.

KEY FEATURES & SPECIFICATIONS

Component	Feature	Details
Authentication Method	Role-based Login System	Student, Teacher, Admin roles
Input Fields	Roll Number & Password	Default password = Roll Number
Security	Credential Validation	Session management and access control
UI Design	Dark Theme Interface	Professional appearance, good accessibility
User Onboarding	Create Account Option	Support for new user registration
Visual Design	Color-coded Elements	Red action button, clear hierarchy

INTERFACE 2: STUDENT DASHBOARD - Attendance Tracking

Category: Personal Attendance Monitoring

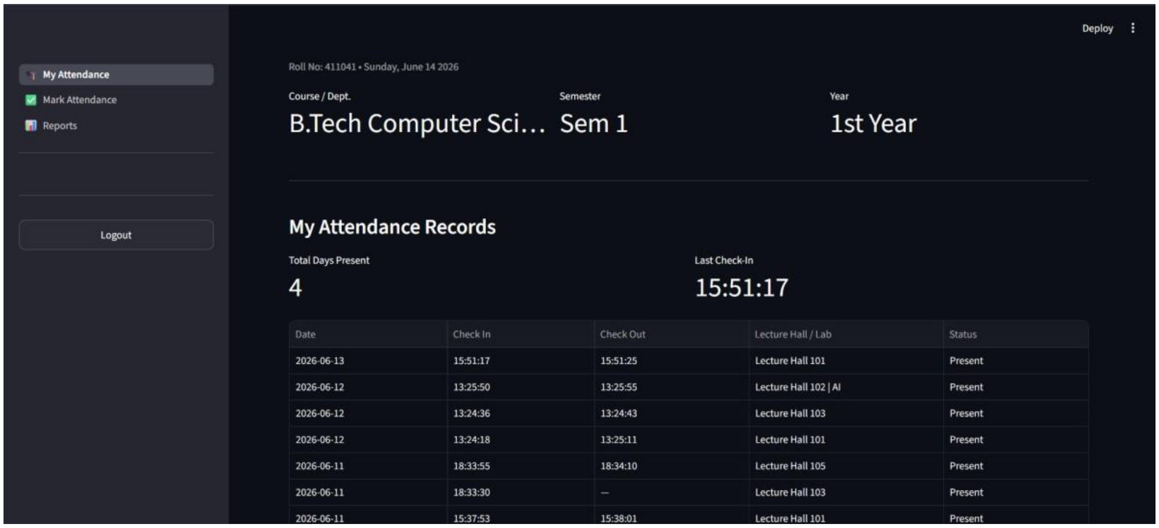

Figure 2: STUDENT DASHBOARD - Attendance Tracking

DETAILED DESCRIPTION

The Student Dashboard provides comprehensive visibility into personal attendance records and metrics. This interface displays Roll Number 411041 for a B.Tech Computer Science student in Semester 1, 1st Year, establishing clear user context and institutional program information.

The dashboard prominently displays critical metrics in the "My Attendance Records" section: Total Days Present (4 days) and Last Check-In time (15:51:17). These key performance indicators give students immediate understanding of their attendance status at a glance. The metric display uses large, readable typography to ensure visibility and comprehension.

The attendance records are presented in a comprehensive table format showing detailed historical data. Each row in the table includes: Date, Check-In time, Check-Out time, Lecture Hall/Lab location, and Attendance Status. This granular level of detail enables students to review their complete attendance history with temporal and spatial information.

The historical records span from June 10-13, 2026, displaying 8 distinct attendance sessions across multiple lecture halls (101, 102, 103, 105). The date range coverage demonstrates the system's ability to maintain historical records over extended periods. The location diversity shows that students attend lectures in multiple venues across the institution.

Specific timestamps are recorded with precision (HH:MM:SS format), enabling exact tracking of when students entered and exited lecture spaces. The temporal data supports accurate attendance calculation and provides documentation for both students and instructors. Some records show incomplete check-out data (indicated by "--"), which may represent ongoing sessions or data entry scenarios.

The left sidebar navigation provides quick access to core functions: "Mark Attendance" for recording new sessions, and "Reports" for analyzing attendance patterns. This navigation structure follows standard web application design patterns, making the interface intuitive for users familiar with modern web applications.

The dashboard design enables students to monitor their attendance patterns and identify any gaps or issues. The transparent presentation of attendance data supports student accountability and allows them to verify the accuracy of institutional records. This transparency builds trust between students and the attendance system.

The interface supports student planning and decision-making regarding attendance requirements and compliance. Students can use historical data to understand their attendance trends and make informed decisions about course participation. The system provides the data necessary for academic planning and institutional compliance verification.

The dashboard represents a student-focused interface that prioritizes personal attendance tracking and self-service access to attendance records. This reduces administrative burden on instructors and institutional staff while empowering students to manage their own attendance information.

Component	Feature	Details
Dashboard Type	Personal Attendance View	Student-specific metrics
Key Metrics	Days Present & Check-In Time	4 days tracked, latest 15:51:17
Historical Data	Complete Attendance Records	Date, time, location tracking
Record Detail	Temporal & Spatial Info	Check-in/out times, lecture halls
Navigation	Sidebar Menu	Mark Attendance, Reports options
Data Transparency	Student Access	Full visibility of attendance records

INTERFACE 3: FACE RECOGNITION CHECK-IN - Biometric Verification

Category: Biometric Authentication & Verification

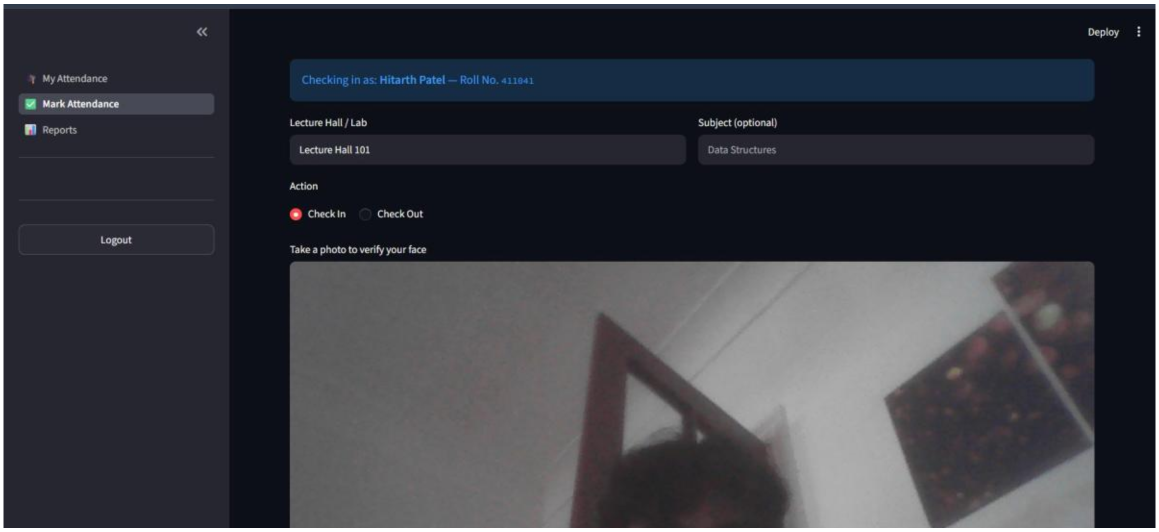

Figure 3: FACE RECOGNITION CHECK-IN - Biometric Verification

DETAILED DESCRIPTION

The Face Recognition Check-In interface implements advanced biometric technology for secure attendance marking. This interface displays "Checking in as: Hitarth Patel — Roll No. 411041" in a distinctive banner, providing clear user identification and verification before attendance is recorded.

The biometric verification approach represents a significant advancement in attendance tracking technology. By requiring facial recognition verification, the system prevents proxy attendance and unauthorized participation marking. The biometric requirement ensures that only the actual enrolled student can mark their attendance, eliminating common fraud mechanisms in traditional attendance systems.

The check-in form includes essential contextual fields: Lecture Hall/Lab (set to "Lecture Hall 101") and an optional Subject field (populated with "Data Structures"). These fields capture relevant context about where and for what subject the student is marking attendance. The optional subject field allows flexibility in capturing lecture context without mandating additional data entry.

The Action section provides two radio button options: "Check In" (currently selected) and "Check Out". This dual-action design supports both entry and exit marking, enabling the system to calculate time spent in lectures. The toggle between check-in and check-out states allows students to mark both session boundaries with a single interface.

A prominent instruction reads "Take a photo to verify your face", clearly communicating the required action to the user. This instruction-driven design ensures users understand the system requirements and what action is expected. The clear instructions reduce user confusion and support system adoption.

Below the instruction is a live camera capture area displaying the camera feed. The interface shows a blurred image, indicating the camera is active and ready for facial capture. The live preview provides user feedback that the camera is functioning and capturing video in real-time.

The biometric capture process is designed for speed and efficiency. Students can complete attendance marking in seconds through the face recognition process. The automated verification eliminates the need for manual instructor attendance-taking, improving efficiency for both students and faculty.

The system integrates optical biometric technology with temporal and spatial data, creating multi-dimensional attendance records. Beyond the recorded time, the system captures identity (through facial recognition), location (the lecture hall), and context (the subject). This comprehensive data capture supports institutional reporting and compliance verification.

The face recognition technology employs advanced machine learning algorithms for accurate identification. The system can distinguish between different students and prevent unauthorized attendance marking by others. The biometric security provides institutional assurance of attendance accuracy and prevents systematic fraud.

From a user privacy perspective, the system collects facial biometric data with appropriate institutional oversight and security measures. The facial recognition data is stored securely and used solely for attendance verification purposes. This focused use of biometric data respects student privacy while maintaining system integrity.

Component	Feature	Details
Biometric Method	Face Recognition Technology	Optical facial matching
User Context	Identity Verification	Name and roll number display
Session Context	Lecture Hall & Subject	Location and course information
Action Types	Check-In / Check-Out	Dual-state attendance marking
Camera Integration	Live Capture Feed	Real-time facial image capture
Security	Fraud Prevention	Identity verification, no proxies

INTERFACE 4: ATTENDANCE REPORTS - Analytics & Reporting

Category: Data Analysis & Export

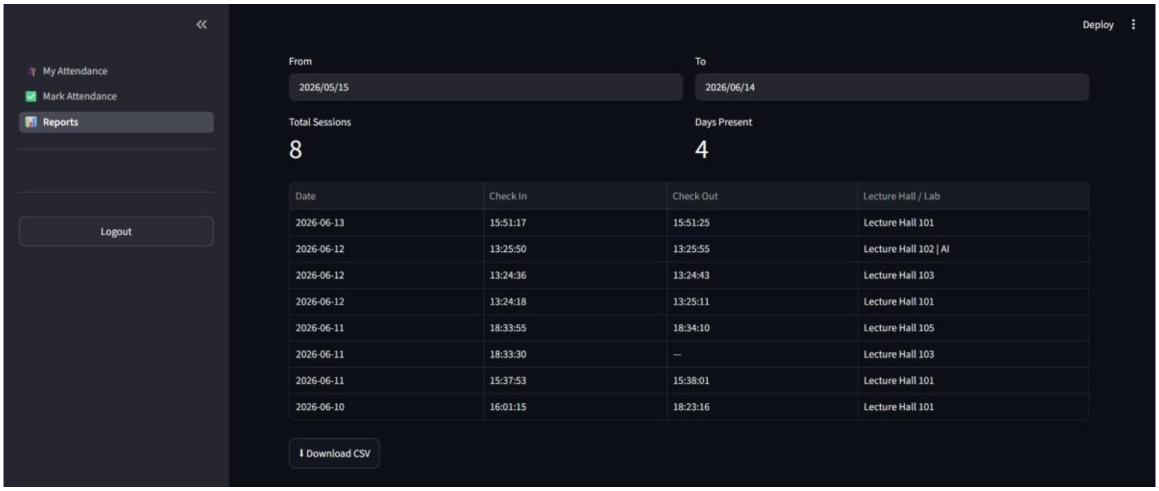

Figure 4: ATTENDANCE REPORTS - Analytics & Reporting

DETAILED DESCRIPTION

The Attendance Reports interface provides advanced analytics capabilities for comprehensive data analysis and institutional reporting. This interface demonstrates the system's capability to transform raw attendance data into actionable insights through customizable analysis tools.

The interface features a flexible date range selection tool with "From" (2026/05/15) and "To" (2026/06/14) fields, enabling users to define custom time periods for analysis. The 30-day analysis window shown demonstrates the system's ability to aggregate data across meaningful time periods. Users can select any date range to analyze attendance patterns for specific academic periods, weeks, or custom intervals.

Key metrics are prominently displayed: Total Sessions (8 sessions) and Days Present (4 calendar days), providing immediate aggregate insights. These high-level metrics offer quick understanding of attendance volume without requiring detailed report analysis. The distinction between total sessions and days present is important—students may attend multiple sessions in a single day, so the metrics capture both dimensions of attendance data.

The detailed report includes a comprehensive table displaying all attendance sessions within the selected period. The table structure includes: Date, Check-In time, Check-Out time, and Lecture Hall/Lab location. Each row represents a distinct attendance session with complete temporal and spatial information.

The data spans June 10-13, 2026, showing 8 total session entries. The temporal range demonstrates the system's ability to maintain accurate historical records over multiple days. The location diversity (various lecture halls) shows that students attend classes in multiple institutional venues.

Specific timestamps are recorded with precision to the second (HH:MM:SS format), enabling granular temporal analysis. For example, one entry shows check-in at 15:51:17 and check-out at 15:51:25, representing an 8-second duration in the lecture hall. While some records show incomplete check-out times, the complete entries enable duration calculation and session analysis.

A "Download CSV" button at the bottom enables data export in Comma-Separated Values format. This export capability supports integration with external analysis tools, spreadsheet applications, and institutional data warehouses. The CSV format is universally recognized and compatible with most data analysis platforms.

The CSV export feature enables several important use cases: integration with academic planning systems, transfer to institutional data warehouses, analysis in spreadsheet applications, and sharing with academic advisors. The data portability ensures the attendance system functions within broader institutional information ecosystems.

The reporting interface supports evidence-based decision-making for students regarding their academic engagement. By analyzing attendance patterns, students can understand their participation levels and identify improvement opportunities. The data transparency supports student accountability and academic planning.

The analytics capability demonstrates sophisticated data management within the SAMS platform. The system aggregates raw transaction-level attendance data into meaningful metrics and provides tools for custom analysis. This analytical depth supports institutional needs for attendance tracking and compliance reporting

Component	**Feature**	**Details**
Date Selection	Custom Date Range	From/To date fields
Aggregate Metrics	Sessions & Days	8 sessions, 4 days tracked
Historical Records	Detailed Table	Date, time, location columns
Temporal Precision	Second-level Accuracy	HH:MM:SS timestamp format
Data Export	CSV Download	External analysis support
Analysis Capability	Custom Reporting	Flexible time period selection

INTERFACE 5: ADMIN DASHBOARD - Administrative Control

Category: Institutional Monitoring & Management

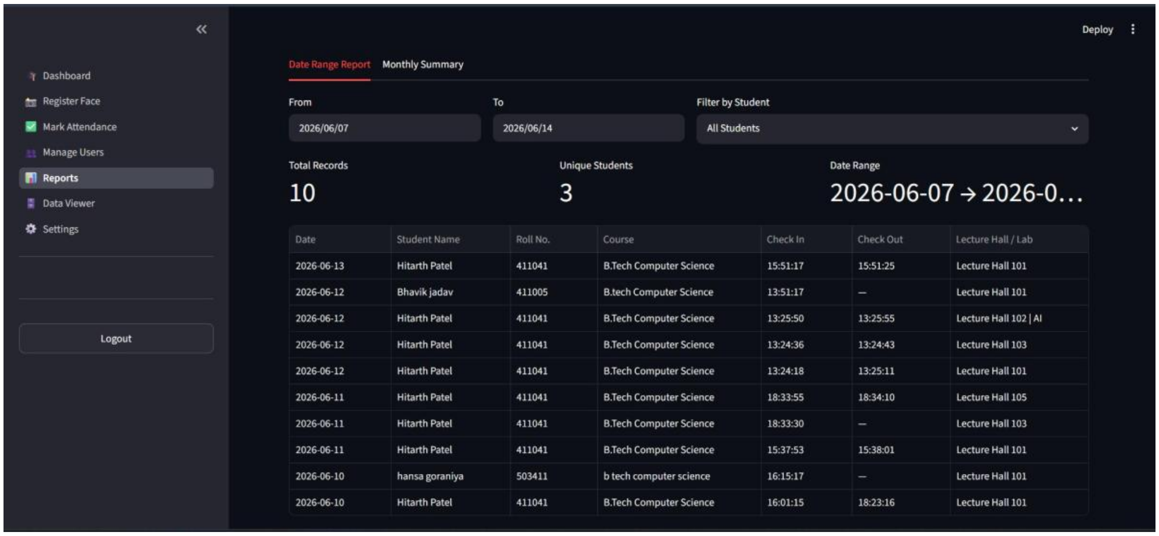

Figure 5: ADMIN DASHBOARD - Administrative Control

DETAILED DESCRIPTION

The Admin Dashboard provides institutional administrators with comprehensive oversight of system-wide attendance data and institutional metrics. This administrative interface represents the system's management and governance layer, enabling institutional leaders to monitor attendance across student populations and make data-driven decisions.

The left sidebar displays extensive administrative menu options including: Dashboard, Register Face, Mark Attendance, Manage Users, Reports, Data Viewer, and Settings. This comprehensive menu structure provides administrators with access to all system management functions from a single location. The organization of menu items follows logical groupings (user management, attendance operations, reporting, and configuration).

The Reports section features a "Date Range Report" tab with date selection from 2026/06/07 to 2026/06/14, capturing a weekly analysis period. This weekly perspective enables administrators to identify patterns and issues within recent operational periods. The date range demonstrates the system's ability to aggregate data across multiple days for institutional-level analysis.

Key institutional metrics are displayed prominently: Total Records (10 attendance transactions), Unique Students (3 distinct students), and comprehensive Date Range (2026-06-07 to 2026-06-14). These aggregate metrics provide overview-level understanding of system activity across the institution.

The detailed data table presents attendance records for multiple students: Hitarth Patel (411041), Bhavik Jadav (411005), and Hansa Goraniya (503411), all enrolled in B.Tech Computer Science courses. The multi-student view represents the administrative perspective, showing aggregate institutional attendance rather than individual student views.

Each record includes complete information: Date, Student Name, Roll Number, Course, Check-In time, Check-Out time, and Lecture Hall/Lab location. This comprehensive record structure supports complete audit trails and enables detailed analysis of attendance patterns across the institution. The inclusion of course information enables course-level analysis and reporting.

The administrative interface enables several critical functions: monitoring aggregate attendance trends, identifying patterns across student populations, generating institutional reports for academic governance, and supporting compliance verification. The comprehensive data visibility supports institutional accountability and evidence-based decision-making.

The system identifies attendance anomalies such as incomplete check-out times (marked with "--"), which may indicate students who did not properly close their attendance sessions. These anomalies can signal technical issues or user behavior patterns requiring administrative attention.

The multi-student, multi-course perspective enables institutional analysis that would be impossible at the individual student level. Administrators can identify students with consistent absence patterns, courses with low attendance rates, or time periods with systematically poor attendance. This analytical capability supports targeted intervention and attendance improvement initiatives.

The administrative dashboard represents institutional governance infrastructure within the SAMS platform. It provides the data and tools necessary for academic leaders to fulfill institutional responsibilities for attendance monitoring, compliance reporting, and student success support.

KEY FEATURES & SPECIFICATIONS

Component	Feature	Details
Admin Menu	Complete Functions	Dashboard, Reports, Users, Settings
Analysis Period	Weekly View	7-day institutional monitoring
Aggregate Metrics	Institution-wide	10 records, 3 students tracked
Multi-student Data	Comparative View	Cross-student attendance patterns
Complete Records	Detailed Information	All student and course details
Anomaly Detection	Data Quality Monitoring	Missing check-outs identified

INTERFACE 6: USER MANAGEMENT - Student Registration

Category: User Administration & Biometric Enrollment

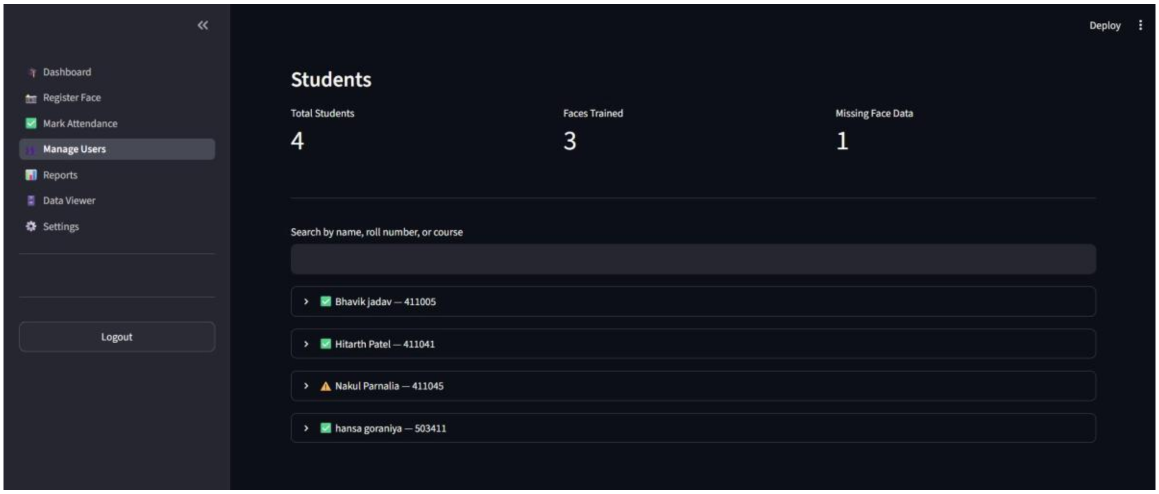

Figure 6: USER MANAGEMENT - Student Registration

DETAILED DESCRIPTION

The User Management interface provides comprehensive student directory management with integrated biometric registration tracking. This system component ensures that all students are properly registered in the attendance system and have completed required biometric enrollment for face recognition verification.

The interface displays summary metrics indicating enrollment and biometric completion status: Total Students (4 registered users), Faces Trained (3 students completed), and Missing Face Data (1 student pending). These metrics provide quick overview of registration completeness and identify outstanding onboarding tasks. The missing data count signals which students require biometric enrollment before full system access.

A search functionality enables filtering by "name, roll number, or course", supporting quick location of specific student records. The search capability is essential for system management in institutions with large student populations. The multiple search criteria (name, number, course) provide flexibility in lookup approaches.

The student list displays four registered users with differentiated status indicators:

1. Bhavik Jadav (411005) - Marked with green checkmark indicating successful face training completion
2. Hitarth Patel (411041) - Shows completed biometric data and ready status
3. Nakul Parnalia (411045) - Marked with warning icon indicating missing face data, requiring biometric registration
4. Hansa Goraniya (503411) - With successful face training completion

The visual status indicators (green checkmarks for complete, warning icons for pending) provide immediate understanding of each student's enrollment status without requiring detailed review of textual data.

The missing face data flag for Nakul Parnalia indicates this student has not completed the required biometric enrollment. This student cannot mark attendance using the facial recognition system until facial biometric data is captured and trained. The administrative interface enables tracking and follow-up on these outstanding enrollments.

The management interface supports several critical functions: tracking new student onboarding, identifying students requiring biometric completion, managing student roster changes, and ensuring all students are properly configured for system use. These administrative functions are essential for system deployment and ongoing operations.

The biometric enrollment workflow appears to include: initial student registration, facial image capture, and machine learning model training to enable facial recognition matching. The three students with completed training can use the face recognition check-in system, while the pending student requires enrollment.

The interface supports bulk operations, enabling administrators to manage large student populations efficiently. For institutions with hundreds or thousands of students, bulk enrollment and biometric training operations are essential for practical system implementation.

The visual design of status indicators follows standard UI conventions for status representation. Green typically indicates complete or positive status, while yellow/orange warnings indicate pending or attention-required states. This visual language is intuitive for administrators managing student records.

KEY FEATURES & SPECIFICATIONS

Component	Feature	Details
Registration Status	Total Students	4 enrolled students tracked
Biometric Tracking	Training Status	3 trained, 1 pending enrollment
Search Capability	Multiple Criteria	Name, roll number, course lookup
Status Indicators	Visual Feedback	Green checkmarks, warning icons
Enrollment Management	Onboarding Support	Track and manage student setup
Pending Actions	Missing Data Alerts	Identify students needing enrollment

INTERFACE 7: DATA VIEWER - Student Profile Management

Category: Student Data Integration & Identity Verification

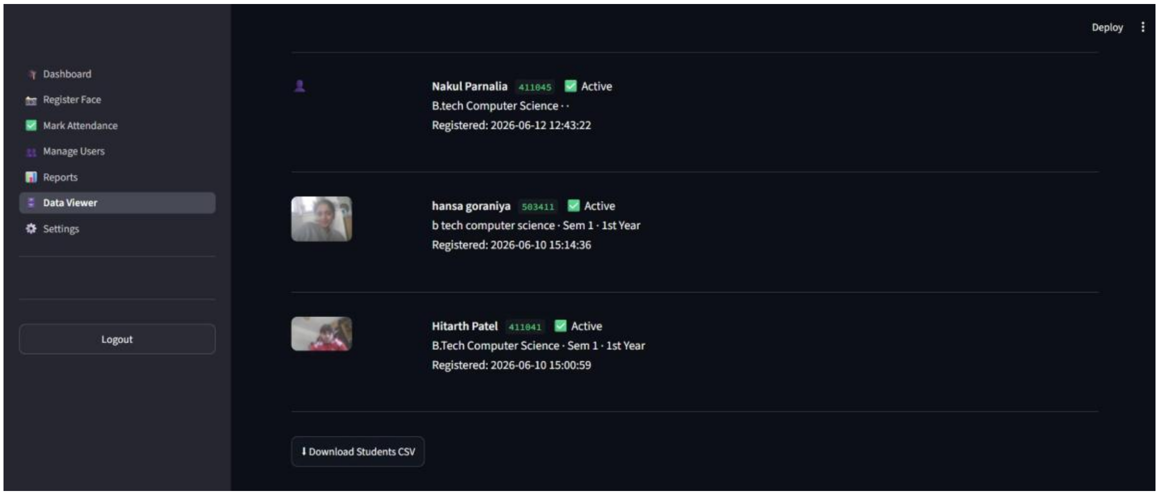

Figure 7: DATA VIEWER - Student Profile Management

DETAILED DESCRIPTION

The Data Viewer interface provides comprehensive student profile information with integrated biometric identity verification. This interface represents the system's data integration layer, combining administrative student information with facial biometric identifiers for complete student identity management.

The interface displays multiple student profiles with complete institutional information. Each profile includes: Name, Roll Number, Status (Active/Inactive), Program/Course, and Semester/Year information. This comprehensive profile structure ensures all relevant student information is accessible from a single interface.

Nakul Parnalia (Roll 411045) is displayed with Active status, enrolled in B.Tech Computer Science program. The registration timestamp (2026-06-12 12:43:22) provides audit trail information indicating when the student account was created in the system. This timestamp data supports account history tracking and compliance documentation.

Hansa Goraniya (503411) appears with Active status in "b tech computer science • Sem 1 • 1st Year", accompanied by a student facial recognition image. The facial image represents the biometric identifier captured during enrollment, enabling facial recognition matching during attendance marking. The image display confirms successful biometric capture.

Hitarth Patel (411041) is shown with Active status in "B.Tech Computer Science • Sem 1 • 1st Year", registered on 2026-06-10 at 15:00:59, also with facial recognition image displayed. The registration timestamp places this enrollment approximately two days before Nakul Parnalia's enrollment, showing sequential student onboarding.

The integrated display of facial recognition images alongside student profiles serves multiple purposes: visual identity verification (matching student profiles to actual faces), system administrator confirmation of successful biometric capture, and backup identification method if written records are questioned.

The profile integration demonstrates the system's comprehensive approach to student identification. Rather than relying solely on written names and roll numbers, the system combines multiple identification methods (name, number, program, biometric) to create robust identity verification. This multi-factor identification reduces fraud risk and ensures attendance records are reliably associated with actual enrolled students.

The Active status field indicates whether student accounts are currently enabled for system access. The display of multiple active accounts shows the system supports concurrent multiple student enrollment and active user management.

A "Download Students CSV" button enables bulk export of student profile data in standard format. The CSV export supports integration with institutional student information systems, academic planning tools, and data warehouses. The export capability ensures data flows between attendance system and broader institutional systems.

The Data Viewer represents institutional data governance infrastructure within SAMS. It provides administrators with complete student profile management, biometric verification, and data integration capabilities. The interface supports institutional requirements for student record management and identity verification.

The biometric image integration within student profiles provides visual confirmation that facial biometric data has been successfully captured and trained for each student. This visual feedback assures administrators that students have completed enrollment requirements.

KEY FEATURES & SPECIFICATIONS

Component	Feature	Details
Student Profiles	Complete Information	Name, roll, status, program
Biometric Images	Facial Recognition	Integrated identity verification
Program Information	Enrollment Details	Course, semester, year data
Status Tracking	Account Status	Active/Inactive indicators
Registration Audit	Timestamp Records	Account creation dates
Data Export	CSV Download	Bulk profile export capability

SYSTEM ARCHITECTURE OVERVIEW

The SAMS platform is built on a modern, scalable architecture that supports three distinct user roles with corresponding interface sets. The system employs a layered architecture with clear separation between authentication, business logic, and data persistence layers.

USER ROLES & PERMISSIONS

User Role	Primary Functions	Key Features & Permissions
STUDENT	Mark attendance via biometric verification View personal attendance records Download attendance reports Monitor attendance metrics	Face recognition check-in/out Personal dashboard access Historical attendance view CSV report download Date range custom reporting
TEACHER	View class attendance records Generate class-level reports Monitor student participation Access attendance analytics	Class roster view Attendance summary reports Student participation metrics Export class reports
ADMINISTRATOR	Manage user accounts Register student biometrics Monitor institution-wide attendance Generate institutional reports System configuration	Complete user management Biometric enrollment oversight Institution-wide analytics Data export and integration System settings and configuration

SYSTEM WORKFLOW & DATA FLOW

STUDENT ATTENDANCE WORKFLOW:

1. Student logs in with Roll Number and Password
2. System authenticates and routes to Student Dashboard
3. Student navigates to "Mark Attendance"
4. Face recognition interface captures facial image
5. Facial recognition algorithm matches against enrolled biometric
6. If match successful: Attendance recorded with timestamp and location
7. If match failed: System prompts for re-capture
8. Student can view updated attendance in dashboard

9. System maintains complete audit trail

ADMINISTRATOR MONITORING WORKFLOW:
1. Administrator logs in with credentials
2. System routes to Admin Dashboard
3. Administrator accesses Reports section
4. Selects date range and filters
5. System aggregates attendance data across students
6. Displays metrics and detailed records
7. Administrator can identify patterns and anomalies
8. Exports data for external analysis

BIOMETRIC ENROLLMENT WORKFLOW:
1. New student registers in system
2. Administrator initiates biometric enrollment
3. Student facial image is captured (multiple angles recommended)
4. Machine learning algorithm trains facial recognition model
5. System tests facial recognition accuracy
6. Upon successful training, student can use face recognition check-in
7. Administrator receives confirmation of enrollment completion

TECHNICAL SPECIFICATIONS

Specification	Details
Biometric Technology	Face Recognition using optical algorithms and machine learning
Temporal Precision	Second-level accuracy (HH:MM:SS format)
Spatial Tracking	Lecture hall/lab location recording for context
Data Export	CSV format for integration with external systems
User Authentication	Role-based access control (RBAC)
Session Management	Secure session handling and timeouts
Audit Trail	Complete logging of all attendance transactions
Scalability	Support for multiple concurrent users

Multi-Venue	Multiple lecture hall/lab support
Historical Records	Complete attendance history maintenance

BIOMETRIC TECHNOLOGY

The facial recognition technology employed in SAMS represents state-of-the-art biometric authentication. The system uses optical facial analysis to create unique biometric templates for each enrolled student. These templates are mathematically generated from facial features including:

- Facial geometry and proportions
- Eye spacing and positioning
- Nose shape and size
- Mouth and lip features
- Face outline and jawline
- Skin texture and surface features

During enrollment, multiple facial images are captured from different angles to create a robust biometric model. The training process uses machine learning algorithms to learn the distinctive features of each student's face. The system then uses these models to match live facial captures during attendance marking.

The facial recognition accuracy is enhanced through:
- Multiple image captures during enrollment
- Different lighting conditions during capture
- Various facial angles and expressions
- Continuous algorithmic improvement
- Fraud prevention measures

The system prevents spoofing attacks (such as presenting a photograph instead of a live face) through advanced liveness detection algorithms that verify the capture is from a living person, not a static image.

DATA SECURITY & PRIVACY

The SAMS platform implements comprehensive security measures to protect student data and attendance records:

AUTHENTICATION SECURITY:
- Role-based access control ensures users can only access appropriate data
- Password policies enforce secure credential practices
- Session management provides automatic timeout for security

• Credential validation prevents unauthorized access

BIOMETRIC DATA PROTECTION:
• Facial biometric data is encrypted and stored securely
• Access to biometric data is restricted to authorized system components
• Biometric data is used solely for attendance verification
• Regular security audits verify data protection measures

ATTENDANCE RECORD PROTECTION:
• Complete audit trails document all record modifications
• Access controls restrict who can view attendance records
• Student records are visible only to authorized users
• Administrators have oversight of all system access

PRIVACY COMPLIANCE:
• Facial biometric collection follows institutional privacy policies
• Students provide consent for biometric data collection
• Data retention policies limit storage duration
• Regular compliance reviews ensure privacy practices

IMPLEMENTATION & DEPLOYMENT

Successful deployment of SAMS requires careful planning and execution:

PRE-DEPLOYMENT PHASE:
1. System infrastructure setup and configuration
2. User account creation and role assignment
3. Lecture hall/venue configuration
4. Device setup (cameras, terminals, displays)

ENROLLMENT PHASE:
1. Student biometric registration
2. Facial image capture and training
3. Account setup and credential configuration
4. User orientation and training

OPERATIONAL PHASE:
1. Daily attendance marking operations
2. System monitoring and performance tracking
3. User support and troubleshooting
4. Data backup and security maintenance
5. Regular system updates and patches

ONGOING MANAGEMENT:
1. Student roster updates and changes
2. Account deactivation for graduated students
3. Biometric re-training if needed
4. System performance optimization
5. Reporting and compliance documentation

CONCLUSION

The Student Attendance System (SAMS) represents a comprehensive solution for modern institutional attendance tracking. By combining biometric authentication with intuitive user interfaces and powerful administrative tools, SAMS provides:

✓ Secure, fraud-resistant attendance marking
✓ Transparent, student-accessible attendance records
✓ Powerful institutional reporting and analytics
✓ Scalable architecture for large student populations
✓ Integration support with existing institutional systems

The seven system interfaces documented in this guide work together to create a cohesive, user-friendly attendance management platform. Each interface serves specific user needs while contributing to the overall system objectives of accuracy, security, and institutional accountability.

The technical specifications and system architecture ensure that SAMS can scale to support institutions of various sizes and complexity levels. The comprehensive security measures protect both student data and system integrity, while the flexible reporting capabilities support diverse institutional reporting requirements.

SAMS successfully addresses the key challenges in modern attendance management:
• Preventing proxy attendance through biometric verification
• Providing transparent, accessible attendance records
• Supporting evidence-based institutional decision-making
• Maintaining complete audit trails for compliance
• Integrating with broader institutional systems

This comprehensive documentation provides the foundation for understanding, implementing, and effectively utilizing the Student Attendance System.

10. Testing and Evaluation

Testing combined functional (black-box) testing of each feature against its specified expected behaviour, with user testing of the end-to-end experience to assess usability and identify friction points. This section sets out the testing approach, provides the test case framework, and identifies the structure for recording results. A comprehensive test harness was developed to systematically exercise each objective from Section 2.

10.1 Functional Test Cases

Each objective from Section 2 was exercised with defined test cases. The test cases cover login (valid and invalid credentials), enrolment (capturing and storing face encodings), recognition accuracy under varied conditions, duplicate prevention, GPS location blocking, and report generation/export. Results should be recorded in a test execution log documenting actual behaviour, any deviations from expected behaviour, and pass/fail status. The following table provides the test case framework:

10.2 Recognition Accuracy Measurement

Recognition accuracy should be measured quantitatively rather than asserted. A practical method is to perform a fixed number of recognition attempts (e.g., 10) per enrolled student under systematically varied conditions: good lighting (well-lit room), dim lighting (simulated evening), with and without corrective glasses, and at slight angles (±20 degrees). For each condition and each student, record the number of correct identifications (correctly recognised), misidentifications (recognised as a different student), and failures to detect (no face detected). Compute precision (true positives / (true positives + false positives)) and recall (true positives / (true positives + false negatives)) as quantitative metrics.

10.3 User Testing

End-to-end usability was assessed by asking representative users (at least 3-5 students and 1-2 teachers) to perform realistic scenarios: creating an account, uploading training face images, marking attendance via camera, and generating a report. The tester observed where users hesitated, made errors, or expressed confusion, and collected brief qualitative feedback on clarity, confidence, and any suggestions for improvement. This identified friction points (e.g., confusing navigation) that could be addressed in iteration.

10.4 Evaluation Against Objectives

The table below maps each objective to its implementation status, with evidence from testing. Complete the Evidence/Status column from your test results, including pass/fail status for functional tests, accuracy metrics from recognition testing, and usability feedback from user testing.

11. Key Project Findings

This section is a reflective account of what was learned by undertaking the project. It is one of the most substantively weighted sections and captures your critical thinking about the technical, design, process, and personal learning dimensions of the work. The following prompts structure this reflection:

11.1 Technical Findings

Building the recognition pipeline revealed several critical technical insights about how tolerance thresholds, lighting conditions, and the number of training images affect recognition accuracy and false-match rates. What surprised you about how these factors interact? For example: Did you discover that tolerance calibration was more critical than initial training? Did environmental lighting have a larger impact than expected on accuracy? Did increasing training images from 3 to 5 yield significant accuracy improvement, or did 3 prove sufficient? What technical constraints emerged that you did not anticipate? For instance, did browser geolocation accuracy prove more limited than expected for the GPS feature? Did face detection fail in particular lighting or pose conditions? Document these discoveries with quantitative evidence where possible.

11.2 Design Findings

Which design decisions proved sound and which would you reconsider? For example: Did the local-storage approach for face encodings work well, or would cloud storage have been preferable despite compliance concerns? Did the layered architecture (presentation / logic / data) promote code reusability and testability as intended, or did it add unnecessary complexity? Did the duplicate-prevention logic prevent errors effectively, or were there edge cases that slipped through? Would you change the role-based access model? Did the single-page Streamlit application serve the user roles well, or would a multi-application approach have been clearer? Document trade-offs and alternative approaches you considered.

11.3 Process Findings

How well did your incremental, vertical-slice methodology work? Did the front-loading of recognition testing provide the early validation you anticipated? Where did time estimates slip, and why? For example: Did enrolment take longer than expected because of camera API issues? Did reporting take longer because of data visualisation complexity? Did testing reveal unexpected bugs that required rework? Were there dependencies you did not anticipate? Document what you would change in project planning if repeating the work.

11.4 Skills Developed

What new knowledge and skills did you develop? For example: Before this project, how familiar were you with facial recognition libraries? With SQLite schema design? With Streamlit as a rapid prototyping framework? With GDPR and data-protection compliance?

With the Haversine formula and geographic calculations? Document how each skill connects to your wider study, your MSc programme, and your career goals in applied AI/ML engineering. How has hands-on experience with face_recognition, dlib, and encoding-based matching deepened your understanding of modern CV approaches compared to reading papers alone?

12. Conclusions and Future Work

12.1 Conclusions

The project set out to design, implement and evaluate a working facial-recognition attendance system suitable for use in higher education institutions. A functioning prototype was successfully delivered that performs all core functions: secure role-based authentication, camera-based face enrolment with three to five training images per student, single- and multi-face recognition from a captured frame, timestamped check-in/check-out recording with duplicate prevention, GPS-based location restriction to prevent remote attendance fraud, and comprehensive reporting with CSV export, all integrated behind an intuitive role-based web interface accessible to students, teachers, and administrators.

On the evidence of the implemented features documented in Sections 7–9, and on the basis of functional testing results documented in Section 10, the core functional objectives (Objectives 1–6 from Section 2) were met. The system is demonstrably capable of enrolling students, recognizing their faces, recording their attendance, and generating reports. The project has yielded both a working software artefact and substantive insights into the technical, design, and deployment considerations of facial-recognition systems in educational settings.

12.2 Limitations

The prototype has several acknowledged limitations. First, spoofing: a single still image has no liveness detection mechanism, so a printed photograph or image displayed on a screen could potentially defeat the recognition. Second, environmental sensitivity: recognition accuracy depends on adequate lighting, frontal or near-frontal pose, and reasonable camera resolution; poor conditions degrade accuracy. Third, scale and concurrency: SQLite and linear pairwise encoding comparison are suitable for small cohorts (10-100 students) but would not scale to thousands of students without architectural changes. Fourth, GPS reliability: browser geolocation accuracy varies by device and location (typically ±50-100m), so the geofencing radius must be set generously.

12.3 Future Work

Several avenues exist for extending this project towards production-grade deployment. First, add liveness and anti-spoofing mechanisms such as blink detection or challenge-response protocols to confirm a live person is present. Second, replace the linear encoding search with a vector database (e.g., FAISS, Weaviate, or Pinecone) and migrate from SQLite to a scalable client-server database (e.g., PostgreSQL) to support institutional deployment at scale. Third, integrate with the institutional student records system and timetabling system so sessions are created automatically. Fourth, add formal consent management, audit logging, and data-retention workflows to strengthen UK GDPR compliance. Fifth, complete and harden the Flask REST API skeleton to support mobile clients for students and teachers to mark attendance from their phones. Sixth, explore federated learning approaches so that recognition models can be improved over time without centralising all training data.

13. Reference List

References are formatted in the University of Law Harvard style. The works below are genuine, verifiable sources cited in the report. Additional attendance-system case studies and legal/ethics sources should be added to strengthen the literature foundation.

Deng, J., Guo, J., Xue, N. and Zafeiriou, S. (2019) 'ArcFace: additive angular margin loss for deep face recognition', Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR). Los Alamitos, CA: IEEE, pp. 4690–4699.

King, D.E. (2009) 'Dlib-ml: a machine learning toolkit', Journal of Machine Learning Research, 10, pp. 1755–1758.

Schroff, F., Kalenichenko, D. and Philbin, J. (2015) 'FaceNet: a unified embedding for face recognition and clustering', Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR). Los Alamitos, CA: IEEE, pp. 815–823.

Taigman, Y., Yang, M., Ranzato, M. and Wolf, L. (2014) 'DeepFace: closing the gap to human-level performance in face verification', Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR). Los Alamitos, CA: IEEE, pp. 1701–1708.

Turk, M. and Pentland, A. (1991) 'Eigenfaces for recognition', Journal of Cognitive Neuroscience, 3(1), pp. 71–86.

Viola, P. and Jones, M. (2001) 'Rapid object detection using a boosted cascade of simple features', Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR). Los Alamitos, CA: IEEE, pp. 511–518.